

**W4995 Applied Machine Learning**

# Introduction to Supervised Learning

02/03/20

Andreas C. Müller

# Supervised Learning

$$(x_i, y_i) \propto p(x, y) \text{ i.i.d.}$$

$$x_i \in \mathbb{R}^p$$


$$y_i \in \mathbb{R}$$

$$\text{learn } f(x_i) \approx y_i$$

such that


$$f(x) \approx y$$

# Nearest Neighbors



$$f(x) = y_i, i = \operatorname{argmin}_j ||x_j - x||$$

# Nearest Neighbors



$$f(x) = y_i, i = \operatorname{argmin}_j ||x_j - x||$$

training set

$$X = \begin{pmatrix} 1.1 & 2.2 \\ 6.7 & 0.5 \\ 2.4 & 9.3 \\ 1.5 & 0.0 \\ 0.5 & 3.5 \\ 5.1 & 9.7 \\ 3.7 & 7.8 \end{pmatrix} \quad y = \begin{pmatrix} 0 \\ 1 \\ 1 \\ 0 \\ 1 \\ 0 \\ 0 \end{pmatrix}$$

test set


# KNN with scikit-learn

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y)


from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=1)
knn.fit(X_train, y_train)
print("accuracy: ", knn.score(X_test, y_test))
y_pred = knn.predict(X_test)
```

accuracy: 0.77

# Influence of Number of Neighbors




# Influence of Number of Neighbors







# Influence of $n\_neighbors$




# Model complexity




# Overfitting and Underfitting



# Overfitting and Underfitting



# Overfitting and Underfitting



# Computational Properties

## Naive

- fit: no time
- memory:  $O(n * p)$
- predict:  $O(n * p)$

$n=n_{\text{samples}}$   $p=n_{\text{features}}$

# Computational Properties

## Naive

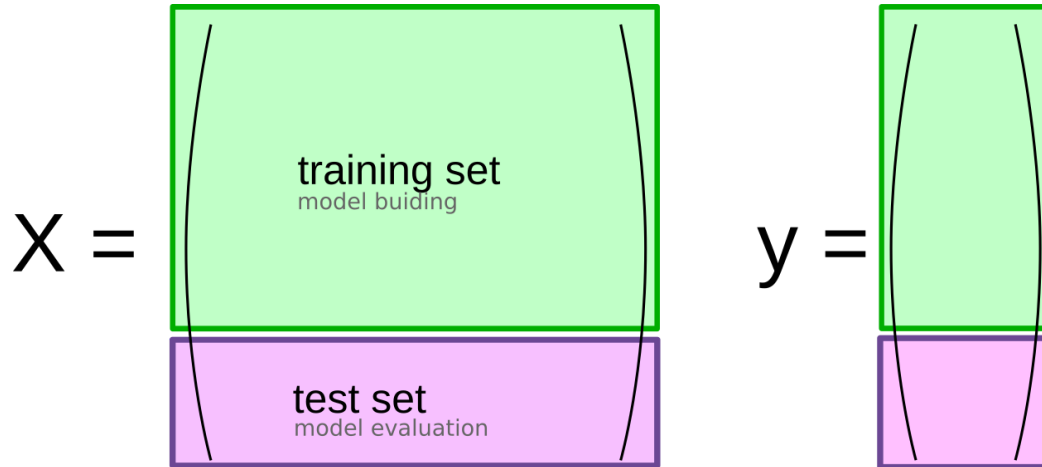
- fit: no time
- memory:  $O(n * p)$
- predict:  $O(n * p)$

## Kd-tree

- fit:  $O(p * n \log n)$
  - memory:  $O(n * p)$
  - predict:
    - $O(k * \log(n))$
- FOR FIXED  $p$ !


$n=n_{\text{samples}}$   $p=p_{\text{features}}$

## So far: Train-test-split

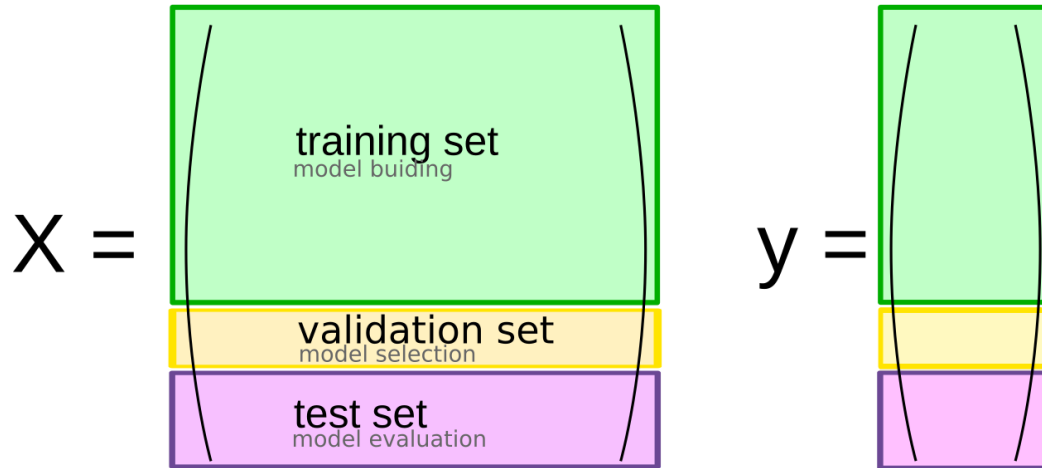





## Threefold split




## Threefold split




# Overfitting the validation set




# Overfitting the validation set



# Overfitting the validation set



# Overfitting the validation set



# Threefold Split for Hyper-Parameters

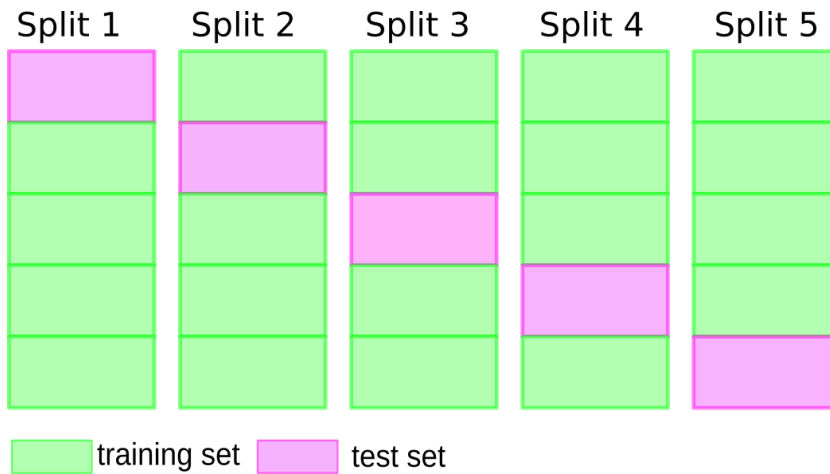
```
X_trainval, X_test, y_trainval, y_test = train_test_split(X, y)
X_train, X_val, y_train, y_val = train_test_split(X_trainval, y_trainval)

val_scores = []
neighbors = np.arange(1, 15, 2)
for i in neighbors:
    knn = KNeighborsClassifier(n_neighbors=i)
    knn.fit(X_train, y_train)
    val_scores.append(knn.score(X_val, y_val))
print(f"best validation score: {np.max(val_scores):.3}")
best_n_neighbors = neighbors[np.argmax(val_scores)]
print("best n_neighbors:", best_n_neighbors)

knn = KNeighborsClassifier(n_neighbors=best_n_neighbors)
knn.fit(X_trainval, y_trainval)
print(f"test-set score: {knn.score(X_test, y_test):.3f}")
```

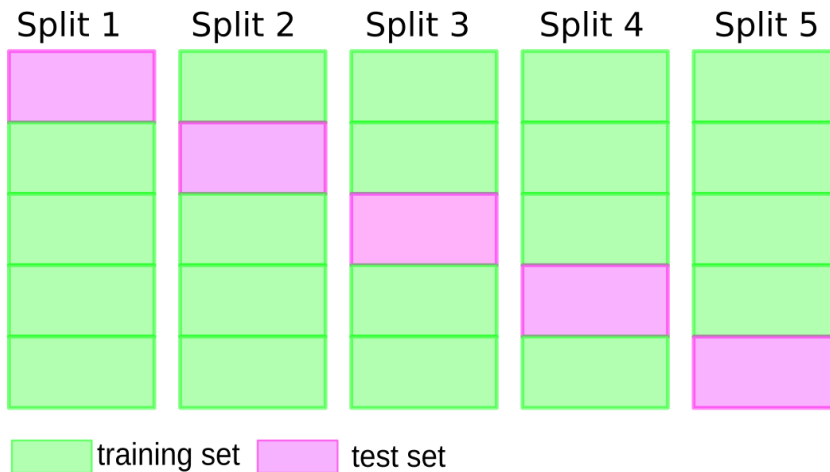
```
best validation score: 0.991
best n_neighbors: 11
test-set score: 0.951
```

# Cross-validation





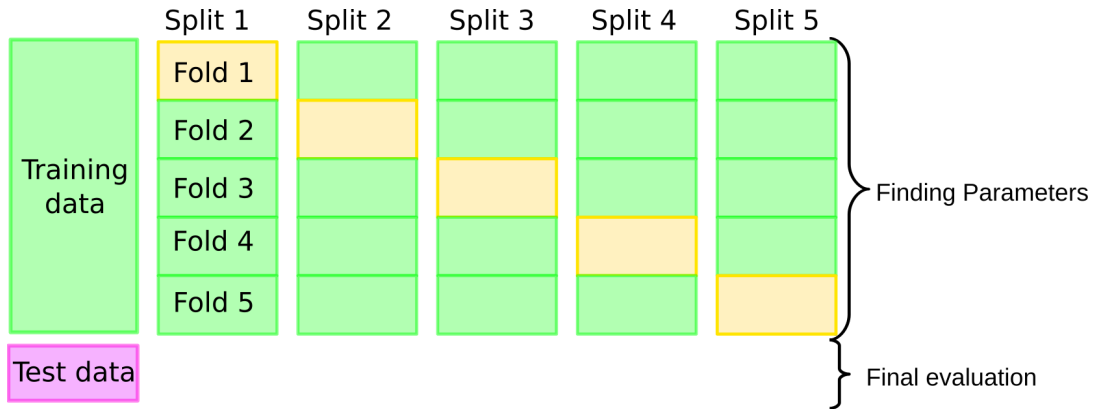
# Cross-validation



pro: more stable, more data

con: slower

# Cross-validation + test set



# Grid-Search with Cross-Validation

```
from sklearn.model_selection import cross_val_score

X_train, X_test, y_train, y_test = train_test_split(X, y)

cross_val_scores = []

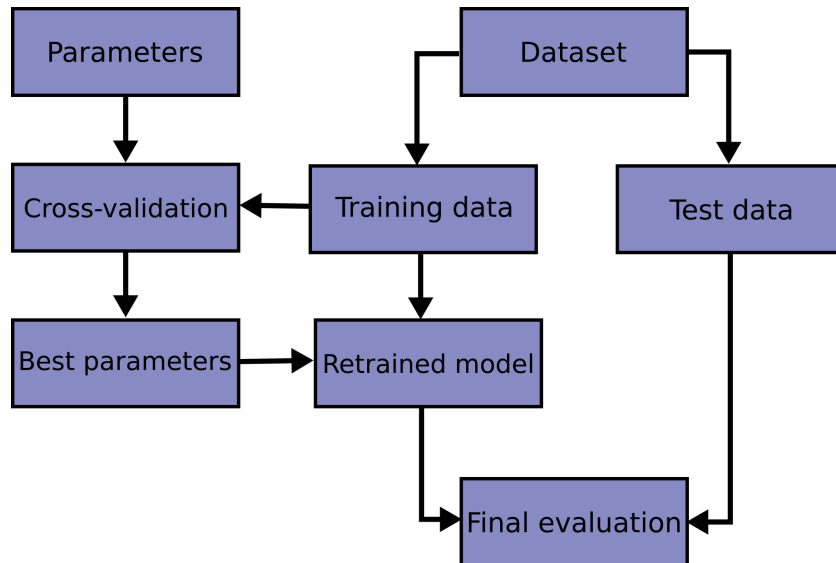
for i in neighbors:
    knn = KNeighborsClassifier(n_neighbors=i)
    scores = cross_val_score(knn, X_train, y_train, cv=10)
    cross_val_scores.append(np.mean(scores))

print(f"best cross-validation score: {np.max(cross_val_scores):.3}")
best_n_neighbors = neighbors[np.argmax(cross_val_scores)]
print(f"best n_neighbors: {best_n_neighbors}")

knn = KNeighborsClassifier(n_neighbors=best_n_neighbors)
knn.fit(X_train, y_train)
print(f"test-set score: {knn.score(X_test, y_test):.3f}")
```

best cross-validation score: 0.967

best n\_neighbors: 9



# GridSearchCV

```
from sklearn.model_selection import GridSearchCV

X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y)

param_grid = {'n_neighbors': np.arange(1, 30, 2)}
grid = GridSearchCV(KNeighborsClassifier(), param_grid=param_grid, cv=10,
                    return_train_score=True)
grid.fit(X_train, y_train)
print(f"best mean cross-validation score: {grid.best_score_}")
print(f"best parameters: {grid.best_params_}")
print(f"test-set score: {grid.score(X_test, y_test):.3f}")
```

```
best mean cross-validation score: 0.967
best parameters: {'n_neighbors': 9}
test-set score: 0.993
```

# GridSearchCV Results

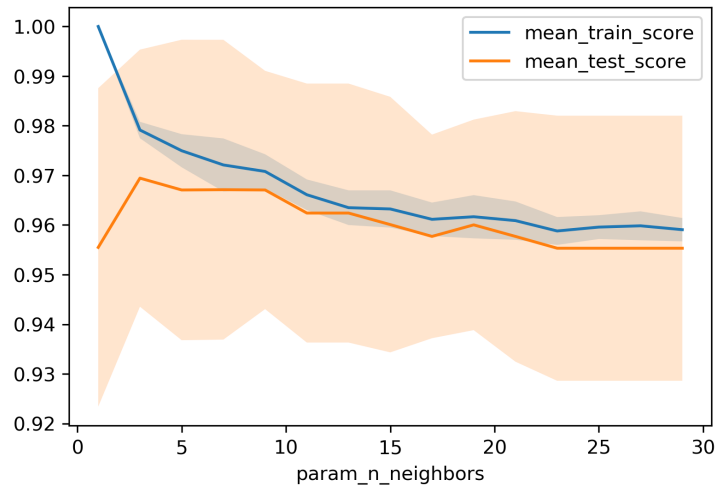
```
import pandas as pd
results = pd.DataFrame(grid.cv_results_)
results.columns
```

```
Index(['mean_fit_time', 'mean_score_time', 'mean_test_score',
      'mean_train_score', 'param_n_neighbors', 'params', 'rank_test_score',
      'split0_test_score', 'split0_train_score', 'split1_test_score',
      'split1_train_score', 'split2_test_score', 'split2_train_score',
      'split3_test_score', 'split3_train_score', 'split4_test_score',
      'split4_train_score', 'split5_test_score', 'split5_train_score',
      'split6_test_score', 'split6_train_score', 'split7_test_score',
      'split7_train_score', 'split8_test_score', 'split8_train_score',
      'split9_test_score', 'split9_train_score', 'std_fit_time',
      'std_score_time', 'std_test_score', 'std_train_score'],
      dtype='object')
```

```
results.params
```

```
0    {'n_neighbors': 1}
1    {'n_neighbors': 3}
2    {'n_neighbors': 5}
3    {'n_neighbors': 7}
4    {'n_neighbors': 9}
5    {'n_neighbors': 11}
6    {'n_neighbors': 13}
Name: params, dtype: object
```

## n\_neighbors Search Results





# Nested Cross-Validation

- Replace outer split by CV loop
- Doesn't yield single model (inner loop might have different best parameter settings)
- Takes a long time, not that useful in practice



# Cross-Validation Strategies





Stratified: Ensure relative class frequencies in each fold reflect relative class frequencies on the whole dataset.

# Importance of Stratification

```
y.value_counts()
```

```
0    60  
1    40
```

```
from sklearn.model_selection import cross_val_score, KFold, StratifiedKFold  
from sklearn.dummy import DummyClassifier
```

```
dc = DummyClassifier('most_frequent')  
skf = StratifiedKFold(n_splits=5, shuffle=True)  
res = cross_val_score(dc, X, y, cv=skf)  
np.mean(res), res.std()
```

```
(0.6, 0.0)
```

```
kf = KFold(n_splits=5, shuffle=True)  
res = cross_val_score(dc, X, y, cv=kf)  
np.mean(res), res.std()
```

```
(0.6, 0.063)
```

# Repeated KFold and LeaveOneOut


- LeaveOneOut : KFold(n\_folds=n\_samples)

High variance, takes a long time


(see [Raschka](#) for a review and [Varoquaux](#) for empirical evaluation)

- Better: ShuffleSplit (aka Monte Carlo)  
Repeatedly sample a test set with replacement

- Even Better: RepeatedKFold.  
Apply KFold or StratifiedKFold multiple times with shuffled data.



Number of iterations and test set size independent



Potentially less variance than StratifiedShuffleSplit.  
Five times five fold or at most ten times ten fold is sufficient.

# Defaults in scikit-learn

- 5-fold in 0.22 (used to be 3 fold)
- For classification cross-validation is stratified
- `train_test_split` has stratify option: `train_test_split(X, y, stratify=y)`
- No shuffle by default!



# Cross-Validation with non-iid data

# Grouped Data


**Assume have data (medical, product, user...) from 5 cities**

- New York, San Francisco, Los Angeles, Chicago, Houston.

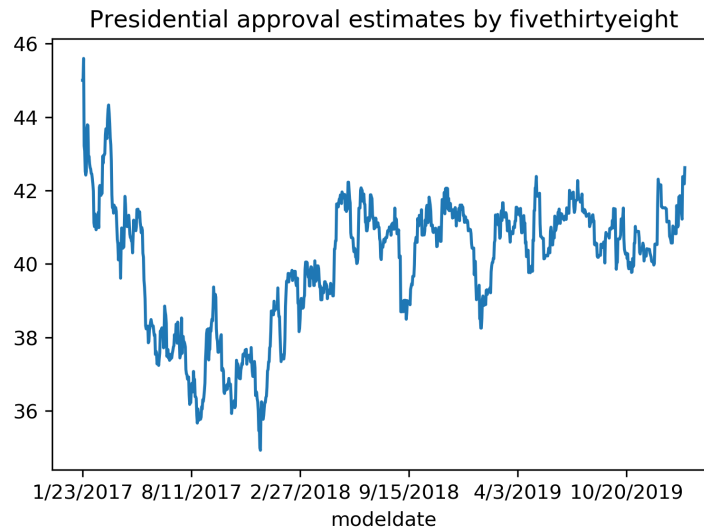
We can assume data within a city is more correlated than between cities.

## **Usage Scenarios**

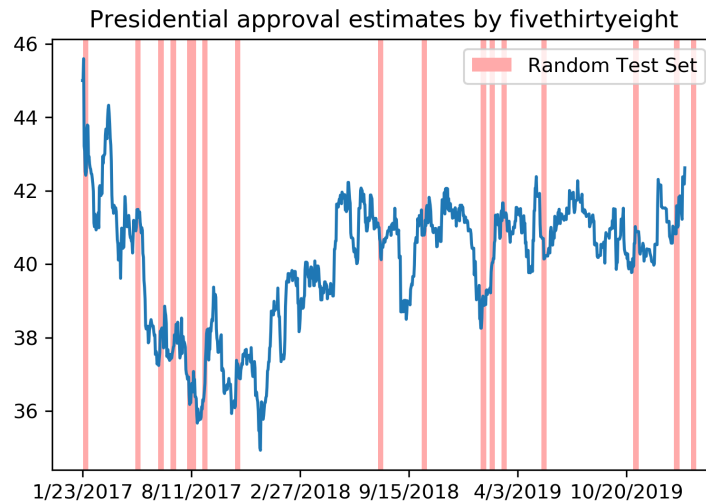
- Assume all future users will be in one of these cities: i.i.d.
- Assume we want to generalize to predict for a new city: not i.i.d.



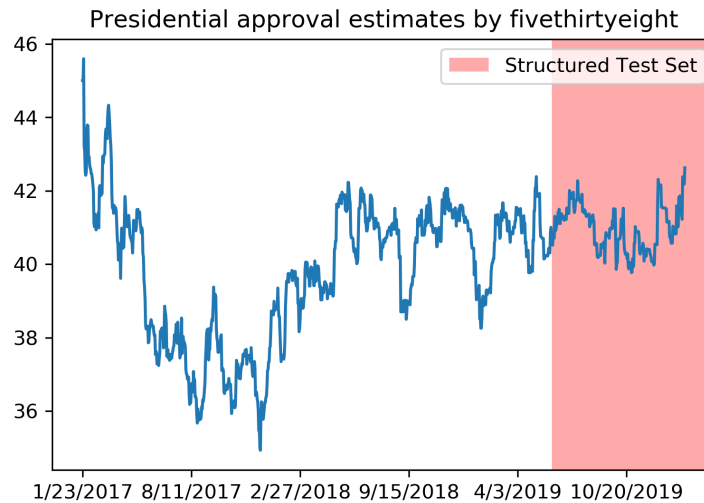
# Correlations in time (and/or space)

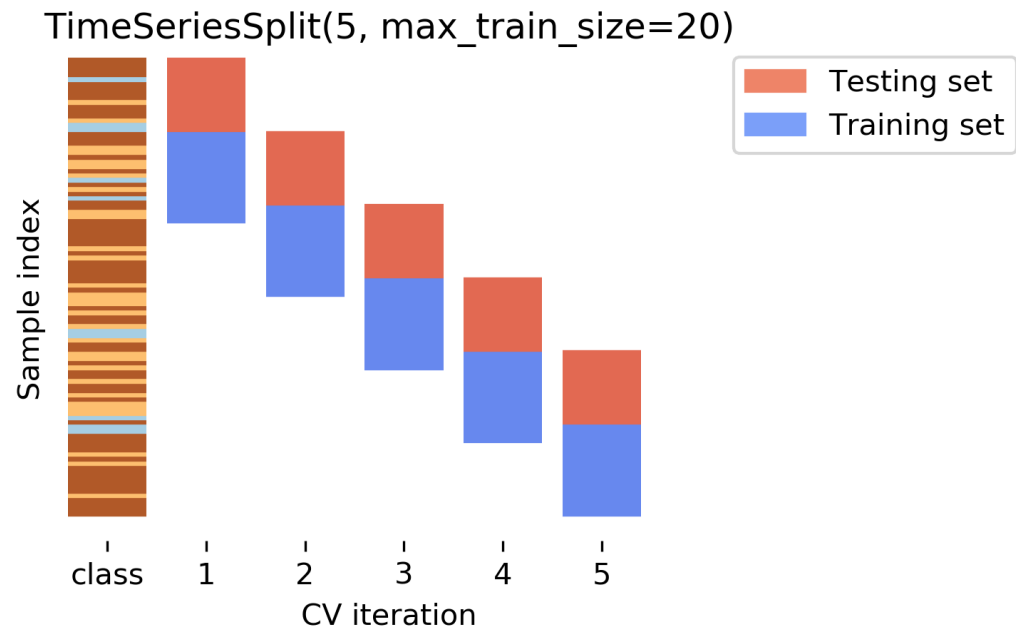



# Correlations in time (and/or space)



# Correlations in time (and/or space)









# Using Cross-Validation Generators

```
from sklearn.model_selection import KFold, StratifiedKFold, ShuffleSplit, RepeatedStratifiedKFold
kfold = KFold(n_splits=5)
skfold = StratifiedKFold(n_splits=5, shuffle=True)
ss = ShuffleSplit(n_splits=20, train_size=.4, test_size=.3)
rs = RepeatedStratifiedKFold(n_splits=5, n_repeats=10)

print("KFold:")
print(cross_val_score(KNeighborsClassifier(), X, y, cv=kfold))

print("StratifiedKFold:")
print(cross_val_score(KNeighborsClassifier(), X, y, cv=skfold))

print("ShuffleSplit:")
print(cross_val_score(KNeighborsClassifier(), X, y, cv=ss))

print("RepeatedStratifiedKFold:")
print(cross_val_score(KNeighborsClassifier(), X, y, cv=rs))
```

```
KFold:
[0.93 0.96 0.96 0.98 0.96]
StratifiedKFold:
[0.98 0.96 0.96 0.97 0.96]
ShuffleSplit:
[0.98 0.96 0.96 0.98 0.94 0.96 0.95 0.98 0.97 0.92 0.94 0.97 0.95 0.92
 0.98 0.98 0.97 0.94 0.97 0.95]
RepeatedStratifiedKFold:
[0.99 0.96 0.97 0.97 0.95 0.98 0.97 0.98 0.97 0.96 0.97 0.99 0.94 0.96
 0.96 0.98 0.97 0.96 0.96 0.97 0.97 0.96 0.96 0.96 0.98 0.96 0.97 0.97
 0.97 0.96 0.96 0.95 0.96 0.99 0.98 0.93 0.96 0.98 0.98 0.96 0.96 0.95
 0.97 0.97 0.96 0.97 0.97 0.97 0.96 0.96]
```

## cross\_validate function

```
from sklearn.model_selection import cross_validate
res = cross_validate(KNeighborsClassifier(), X, y, return_train_score=True,
                    scoring=["accuracy", "roc_auc"])
res_df = pd.DataFrame(res)
```

fit_time	score_time	test_accuracy	test_roc_auc	train_accuracy	train_roc_auc
0.000839	0.010204	0.965217	0.996609	0.980176	0.997654
0.000870	0.014424	0.956522	0.983689	0.975771	0.998650
0.000603	0.009298	0.982301	0.999329	0.971491	0.996977
0.000698	0.006670	0.955752	0.984071	0.978070	0.997820
0.000611	0.006559	0.964602	0.994634	0.978070	0.998026

Questions ?